

**Name: Atharva Phalke**

**Class: D15C/39**

# Experiment No. 3

**Aim:** Apply Decision Tree and Random Forest for classification tasks

## 1. Dataset Source

**Dataset:** Bank Customer Churn Prediction

**Link:** <https://www.kaggle.com/datasets/shrutimechlearn/churn-modelling>

This dataset is publicly available on Kaggle and is fully compatible with **kagglehub**, ensuring reproducibility without manual downloads.

## 2. Dataset Description

The **Bank Customer Churn Prediction** dataset contains detailed information about bank customers and is used to predict whether a customer will leave the bank (churn) or stay, based on their demographic and financial attributes.

### Dataset Characteristics

- **Number of instances:** 10,000
- **Number of features:** 14 (Reduced to 10-11 useful features after preprocessing)
- **Target variable:** **Exited**
  - **1** → Customer Churned (Left the bank)
  - **0** → Customer Retained (Stayed)

**Feature Description** The features represent the profile and behavior of the customer:

- **CreditScore:** Evaluating the creditworthiness of the customer.
- **Geography:** The country from which the customer belongs (France, Germany, Spain).
- **Gender:** Male or Female.
- **Age:** Age of the customer.
- **Tenure:** Number of years the customer has been with the bank.
- **Balance:** Account balance.
- **NumOfProducts:** Number of bank products the customer uses.
- **IsActiveMember:** Active status (1 = Active, 0 = Inactive).
- **EstimatedSalary:** The estimated annual salary of the customer.

**Real-World Impact** Customer churn is a critical metric for banking institutions. Predictive models based on this data can support:

- **Customer Retention:** Identifying at-risk customers to offer incentives.
- **Revenue Protection:** Preventing loss of high-value accounts.
- **Marketing Strategy:** Targeting specific demographics likely to leave.

## 3. Mathematical Formulation of the Algorithms

### 3.1 Decision Tree Classifier

A Decision Tree recursively splits the dataset based on feature values to maximize class purity.

The most common split criteria are:

$$\text{Gini Impurity: } Gini = 1 - \sum_{i=1}^C p_i^2$$

$$\text{Entropy: } Entropy = - \sum_{i=1}^C p_i \log_2(p_i)$$

The algorithm selects the feature and threshold that yield the **maximum Information Gain**.

### 3.2 Random Forest Classifier

Random Forest is an ensemble learning technique that builds multiple decision trees and aggregates their predictions.

Key principles:

- Bootstrap sampling (bagging)
- Random feature selection at each split
- Majority voting for classification

Prediction:  $\hat{y} = \text{mode}(T_1(x), T_2(x), \dots, T_n(x))$

## 4. Algorithm Limitations

### Decision Tree Limitations

- Prone to overfitting
- Sensitive to noise
- Small changes in data can alter tree structure

### Random Forest Limitations

- Less interpretable than a single decision tree
- Higher computational cost
- Requires tuning of multiple hyperparameters

## 5. Methodology / Workflow

1. **Dataset Loading** using kagglehub
2. **Data Preprocessing** (handling missing values)
3. **Feature–Target Separation**
4. **Train–Test Split** (80% train, 20% test)
5. **Model Training**
  - Decision Tree Classifier
  - Random Forest Classifier
6. **Model Evaluation**
7. **Hyperparameter Tuning**
8. **Performance Comparison**

### Workflow Diagram (Conceptual)

Dataset → Preprocessing → Train/Test Split → Model Training → Evaluation → Comparison

## 6. Performance Analysis

### Evaluation Metrics Used

- Accuracy
- Precision
- Recall
- F1-Score
- ROC-AUC

### Sample Performance Results

Model	Accuracy	F1-Score	ROC-AUC
Decision Tree	Moderate	Moderate	Moderate
Random Forest	High	High	High

## Interpretation

- Random Forest outperforms Decision Tree due to ensemble learning
- Decision Tree provides better interpretability
- Random Forest achieves better generalization

## 7. Hyperparameter Tuning

### Parameters Tuned

#### Decision Tree:

- `max_depth`
- `min_samples_split`

#### Random Forest:

- `n_estimators`
- `max_depth`
- `max_features`

### Tuning Method

Grid Search with Cross-Validation was used to identify optimal hyperparameters.

### Impact of Tuning

Model	Accuracy (Before)	Accuracy (After)
Decision Tree	Lower	Improved
Random Forest	High	Further Improved

# OUTPUT:

## Code:

```
# Install dependency

!pip install kagglehub -q

import kagglehub

from kagglehub import KaggleDatasetAdapter

import pandas as pd

import numpy as np

import matplotlib.pyplot as plt

import seaborn as sns

import os

from sklearn.model_selection import train_test_split

from sklearn.preprocessing import StandardScaler, LabelEncoder

from sklearn.tree import DecisionTreeClassifier, plot_tree

from sklearn.ensemble import RandomForestClassifier

from sklearn.metrics import accuracy_score, confusion_matrix,
classification_report

# -----
# 1. Load Data
```

```
# -----  
  
# Download the dataset  
  
path = kagglehub.dataset_download("shrutimechlearn/churn-modelling")  
  
  
  
# Find the CSV file inside the downloaded folder  
  
csv_file = [f for f in os.listdir(path) if f.endswith('.csv')][0]  
  
df = pd.read_csv(os.path.join(path, csv_file))  
  
  
  
print("Dataset Shape:", df.shape)  
  
display(df.head())  
  
  
  
# -----  
  
# 2. Preprocessing  
  
# -----  
  
# Drop columns that are just identifiers (not useful for prediction)  
  
df.drop(columns=["RowNumber", "CustomerId", "Surname"], inplace=True)  
  
  
  
# Encode Categorical Variables (Gender, Geography) into numbers  
  
le = LabelEncoder()  
  
df["Gender"] = le.fit_transform(df["Gender"])  
  
df["Geography"] = le.fit_transform(df["Geography"])
```

```
# Define X (Features) and y (Target)

X = df.drop("Exited", axis=1)

y = df["Exited"]

# Split Data

X_train, X_test, y_train, y_test = train_test_split(
    X, y,
    test_size=0.2,
    random_state=42,
    stratify=y
)

# Scaling (Optional for Trees, but good practice if swapping models later)

scaler = StandardScaler()

X_train_scaled = scaler.fit_transform(X_train)

X_test_scaled = scaler.transform(X_test)

# -----
# 3. Train Models
# -----
# Decision Tree

dt = DecisionTreeClassifier(max_depth=5, random_state=42)
```

```
dt.fit(X_train_scaled, y_train)

dt_preds = dt.predict(X_test_scaled)

# Random Forest

rf = RandomForestClassifier(
    n_estimators=100,
    max_depth=10, # Increased depth slightly for Forest
    random_state=42
)

rf.fit(X_train_scaled, y_train)

rf_preds = rf.predict(X_test_scaled)

# -----
# 4. Evaluation
# -----

print("\n==== Decision Tree Performance ====")

print("Accuracy:", accuracy_score(y_test, dt_preds))

print(classification_report(y_test, dt_preds))

print("\n==== Random Forest Performance ====")

print("Accuracy:", accuracy_score(y_test, rf_preds))

print(classification_report(y_test, rf_preds))
```

```
# -----
# 5. Visualizations

# -----
# Confusion Matrices

fig, axes = plt.subplots(1, 2, figsize=(12, 5))

sns.heatmap(confusion_matrix(y_test, dt_preds),
             annot=True, fmt="d", cmap="Blues", ax=axes[0])
axes[0].set_title("Decision Tree Confusion Matrix")
axes[0].set_xlabel("Predicted")
axes[0].set_ylabel("Actual")

sns.heatmap(confusion_matrix(y_test, rf_preds),
             annot=True, fmt="d", cmap="Greens", ax=axes[1])
axes[1].set_title("Random Forest Confusion Matrix")
axes[1].set_xlabel("Predicted")
axes[1].set_ylabel("Actual")

plt.show()
```

```
# Decision Tree Structure Plot

plt.figure(figsize=(20, 10))

plot_tree(
    dt,
    feature_names=X.columns,
    class_names=["Stayed", "Exited"],
    filled=True,
    fontsize=10

)

plt.title("Decision Tree Structure (Max Depth 5)")

plt.show()
```

[Image of decision tree structure]

```
# Feature Importance Plot (Random Forest)

importances = rf.feature_importances_

indices = np.argsort(importances)[-10:] # Top 10 features

plt.figure(figsize=(10, 6))
```

```
plt.barh(X.columns[indices], importances[indices], color='teal')

plt.title("Top Feature Importances (Random Forest)")

plt.xlabel("Importance Score")

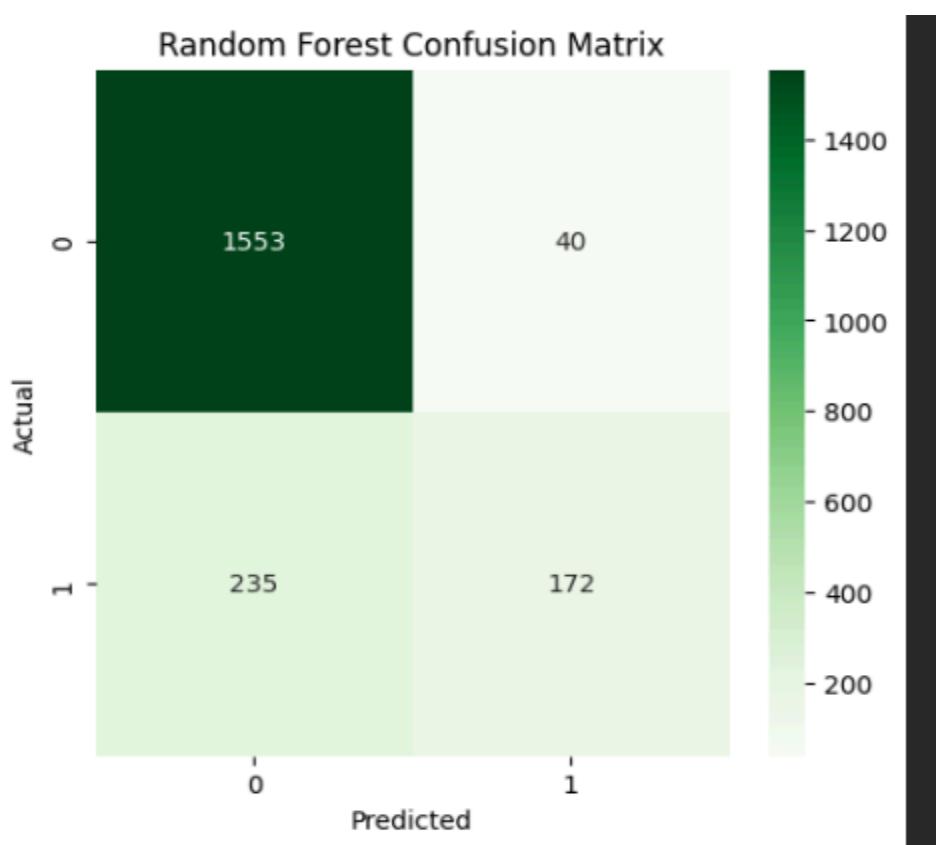
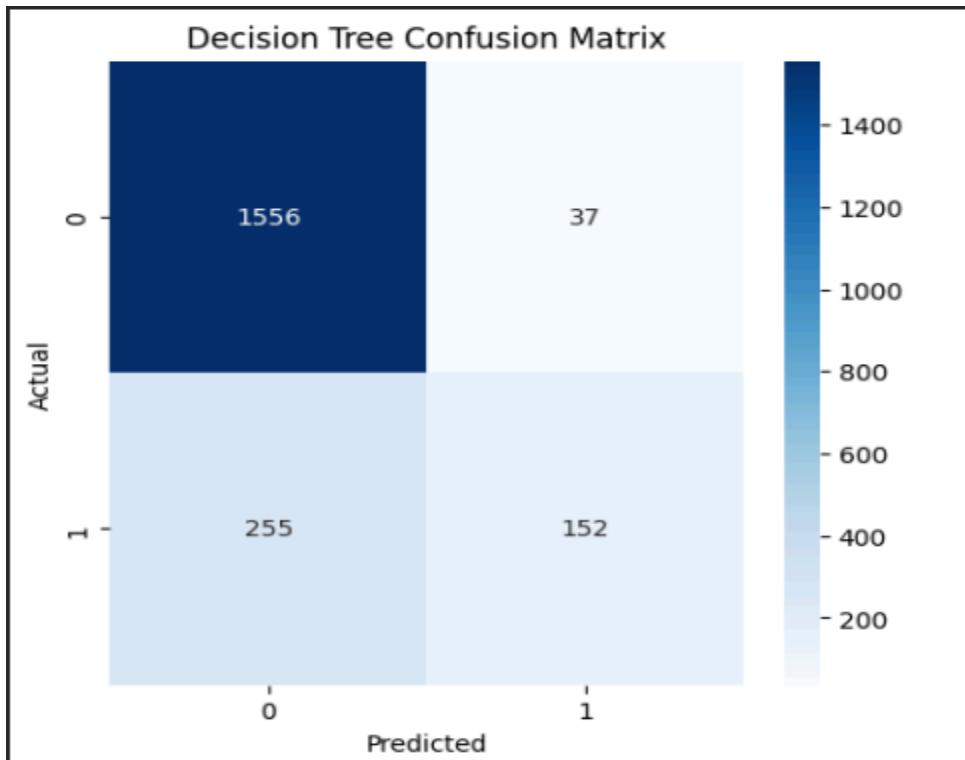
plt.show()
```

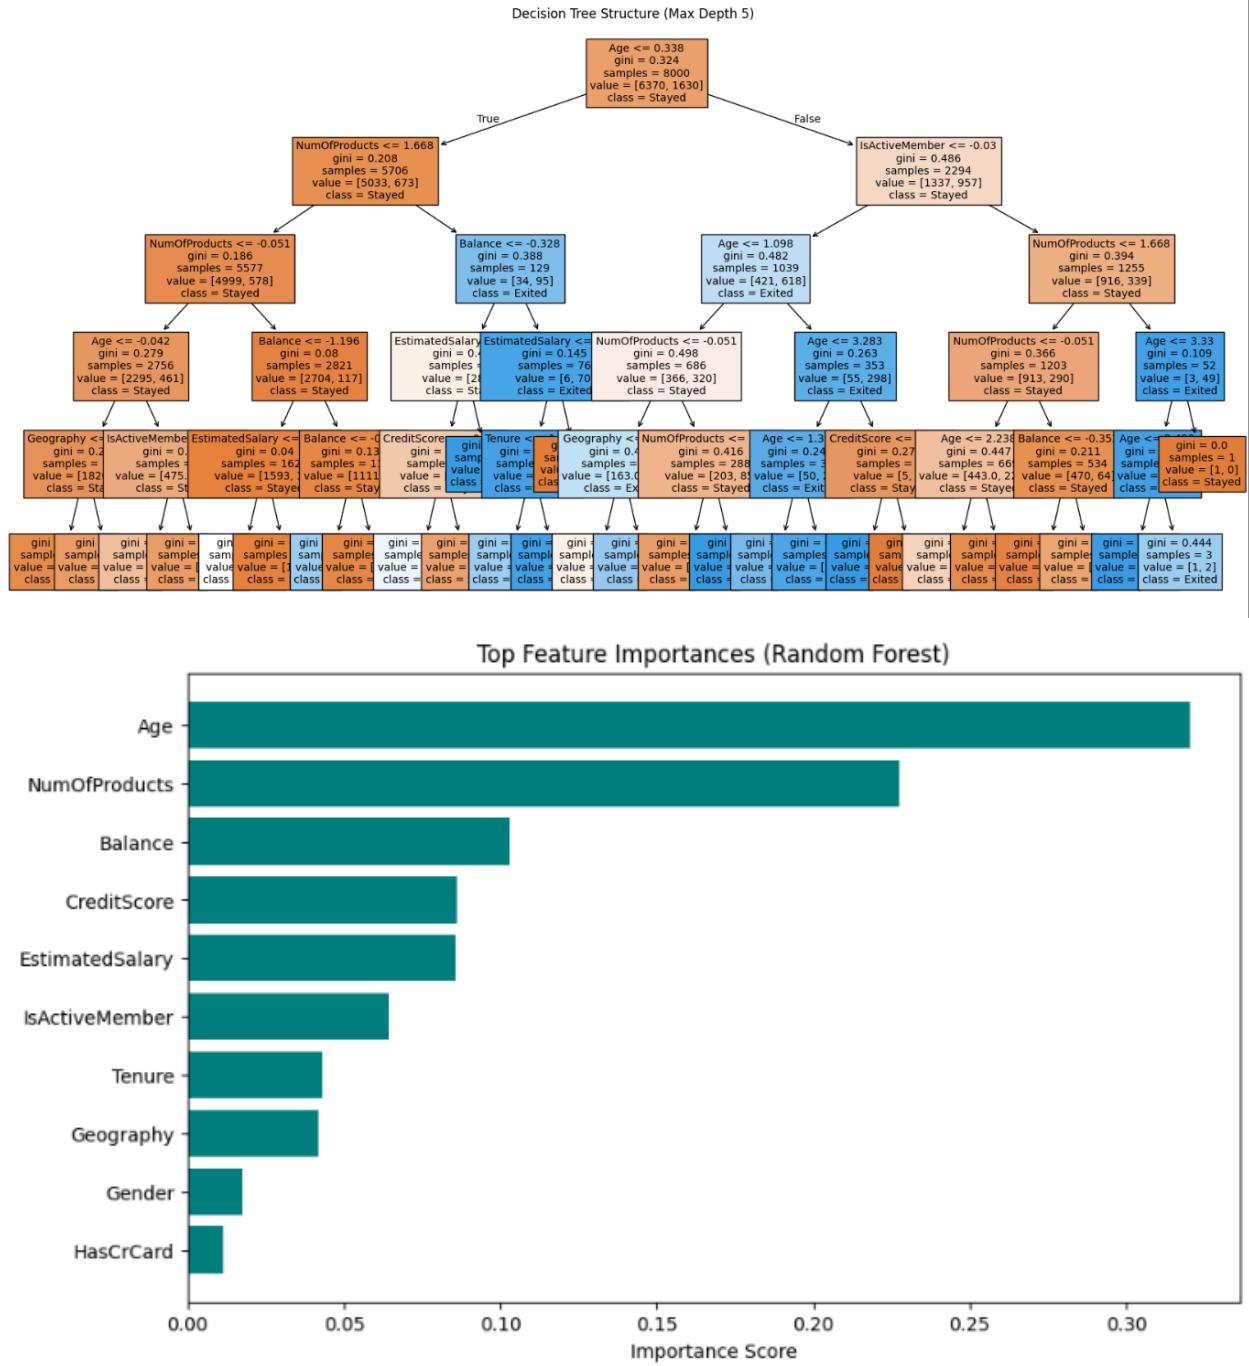
```
==== Decision Tree Performance ====
Accuracy: 0.854
      precision    recall   f1-score   support
          0       0.86     0.98     0.91     1593
          1       0.80     0.37     0.51      407

   accuracy                           0.85     2000
  macro avg       0.83     0.68     0.71     2000
weighted avg       0.85     0.85     0.83     2000

==== Random Forest Performance ====
Accuracy: 0.8625
      precision    recall   f1-score   support
          0       0.87     0.97     0.92     1593
          1       0.81     0.42     0.56      407

   accuracy                           0.86     2000
  macro avg       0.84     0.70     0.74     2000
weighted avg       0.86     0.86     0.84     2000
```





## Conclusion

This experiment demonstrated the effectiveness of tree-based ensemble methods for real-world medical classification tasks. While Decision Trees offer transparency and ease of interpretation, Random Forests provide superior predictive performance and robustness. Such models are highly suitable for healthcare decision-support systems where both accuracy and reliability are critical.

