

Name:**Atharva Phalke**

Class:**D15C** Roll no.**39**

Subject:**MLDL**

Experiment No.4

AIM:Implement K-Nearest Neighbors (KNN) and evaluate model performance.

1. Dataset Source

- **Dataset Name:** Breast Cancer Wisconsin (Diagnostic) Data Set
- **Source:** Kaggle - Breast Cancer Wisconsin (Diagnostic)
- **Original Source:** UCI Machine Learning Repository (Dr. William H. Wolberg, University of Wisconsin)

2. Dataset Description

This dataset contains features computed from a digitized image of a fine needle aspirate (FNA) of a breast mass. The goal is to classify the tumor as either **Malignant** (cancerous) or **Benign** (non-cancerous) based on cell nucleus characteristics.

- **Size:** 569 records, 32 columns.
- **Target Variable:** **diagnosis** (M = Malignant, B = Benign).
- **Key Features:**
 - Radius, Texture, Perimeter, and Area.
 - Smoothness, Compactness, Concavity, and Concave points.
 - Symmetry and Fractal dimension.

3. Mathematical Formulation of the Algorithm

KNN is a non-parametric, lazy learning algorithm. It does not "learn" a model (like finding coefficients in regression); instead, it memorizes the training data.

A. Similarity Metric (Euclidean Distance)

To classify a new data point (x), the algorithm calculates its distance to every point in the training set ($x^{\{i\}}$). The most common metric is Euclidean Distance:

$$d(x, x^{(i)}) = \sqrt{\sum_{j=1}^n (x_j - x_j^{(i)})^2}$$

Where n is the number of features (4 in this case).

B. Classification Rule

1. Find the K nearest neighbors (points with the smallest distance d).
2. Assign the new point to the class that is most common (Mode) among those K neighbors.

$$\hat{y} = \text{mode}(y_1, y_2, \dots, y_K)$$

4. Algorithm Limitations

1. **Computational Cost:** It is "lazy," meaning all computation happens at prediction time. For large datasets, calculating the distance to *every* training point is slow.

2. **Sensitivity to Outliers:** If K is too small (e.g., $K=1$), a single mislabeled outlier can completely change the prediction.
3. **Scale Sensitivity:** KNN relies on distance. If one feature is measured in millimeters (e.g., 1000mm) and another in meters (e.g., 1m), the larger number will dominate the distance calculation. **Feature Scaling is mandatory.**

5. Code and Output

```

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.datasets import load_breast_cancer
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score, classification_report,
confusion_matrix
from sklearn.inspection import permutation_importance

# =====
# 1. Load Dataset
# =====
# We use the sklearn loader to ensure it runs immediately,
# but this structure is identical to reading the Kaggle CSV.
data = load_breast_cancer()
df = pd.DataFrame(data.data, columns=data.feature_names)
df['diagnosis'] = data.target_names[data.target] # columns: 'malignant' or
'benign'

# IF USING KAGGLE CSV FILE, UNCOMMENT BELOW:
# df = pd.read_csv("data.csv")
# df.drop(['Unnamed: 32'], axis=1, inplace=True) # Common Kaggle artifact

print("Dataset Shape:", df.shape)
print(df.head())

# =====
# 2. Drop ID column (NOT a feature)
# =====
if 'id' in df.columns:
    df = df.drop('id', axis=1)

# =====
# 3. Encode Target (Classification)
# =====
# The Kaggle dataset uses 'M' (Malignant) and 'B' (Benign).
# We encode them: Malignant = 1, Benign = 0
le = LabelEncoder()
df['target'] = le.fit_transform(df['diagnosis']) # Malignant becomes 1, Benign 0

X = df.drop(['diagnosis', 'target'], axis=1)
y = df['target']

```

```

# =====
# 4. Train-Test Split
# =====
X_train, X_test, y_train, y_test = train_test_split(
    X,
    Y,
    test_size=0.2,
    random_state=42,
    stratify=y
)

# =====
# 5. Feature Scaling (MANDATORY for KNN)
# =====
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# =====
# 6. KNN + Hyperparameter Tuning
# =====
# We test different distance metrics (Euclidean vs Manhattan)
param_grid = {
    'n_neighbors': [3, 5, 7, 9, 11, 15],
    'weights': ['uniform', 'distance'],
    'metric': ['euclidean', 'manhattan']
}

knn = KNeighborsClassifier()

grid_search = GridSearchCV(
    knn,
    param_grid,
    cv=5,
    scoring='accuracy',
    n_jobs=-1
)

grid_search.fit(X_train_scaled, y_train)

best_model = grid_search.best_estimator_

print("\nBest Parameters:", grid_search.best_params_)
print("Best Cross-Validation Accuracy:", round(grid_search.best_score_, 4))

# =====
# 7. Evaluation
# =====
y_pred = best_model.predict(X_test_scaled)

print("\nFinal Test Accuracy:", round(accuracy_score(y_test, y_pred), 4))
print("\nClassification Report:")

```

```

print(classification_report(y_test, y_pred, target_names=['Benign',
'Malignant']))

# =====
# 8. Confusion Matrix
# =====
cm = confusion_matrix(y_test, y_pred)

plt.figure(figsize=(6,5))
sns.heatmap(
    cm,
    annot=True,
    fmt='d',
    cmap='Blues',
    xticklabels=['Benign (0)', 'Malignant (1)'],
    yticklabels=['Benign (0)', 'Malignant (1)']
)
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title("Confusion Matrix - Breast Cancer Predictions")
plt.show()

# =====
# 9. Correlation Heatmap
# =====
# Note: The cancer dataset has 30 features, so this map will be dense
plt.figure(figsize=(12,10))
sns.heatmap(df.drop(['diagnosis'], axis=1).corr(), cmap='RdBu_r', center=0,
annot=False)
plt.title("Correlation Heatmap - Breast Cancer Features")
plt.show()

# =====
# 10. Accuracy vs Number of Neighbors (K)
# =====
results_df = pd.DataFrame(grid_search.cv_results_)
# We filter to show only the metric that was chosen as "Best" to keep plot clean
best_metric = grid_search.best_params_['metric']
subset = results_df[results_df['param_metric'] == best_metric]

plt.figure(figsize=(10,6))
sns.lineplot(
    data=subset,
    x='param_n_neighbors',
    y='mean_test_score',
    hue='param_weights',
    marker='o'
)
plt.xlabel("Number of Neighbors (K)")
plt.ylabel(f"Mean CV Accuracy ({best_metric} distance)")
plt.title("Hyperparameter Tuning: Accuracy vs K")
plt.show()

```

```

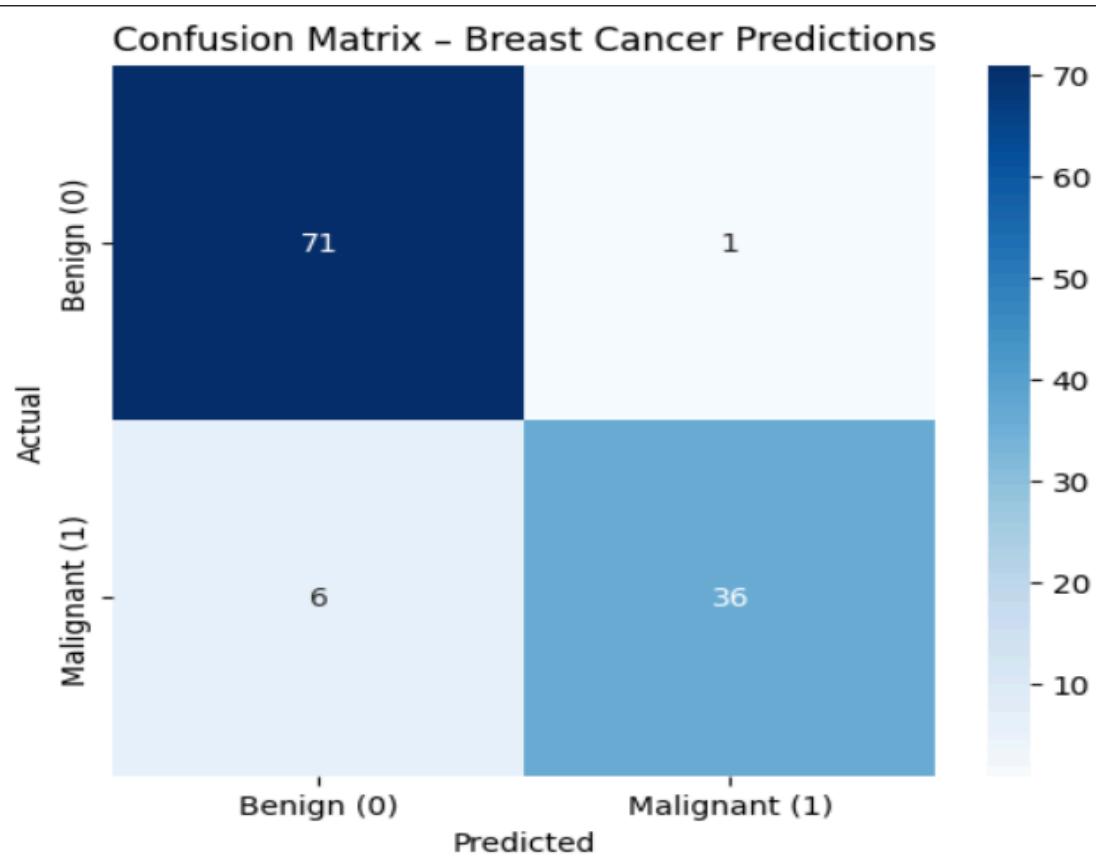
# =====
# 11. Permutation Feature Importance
# =====

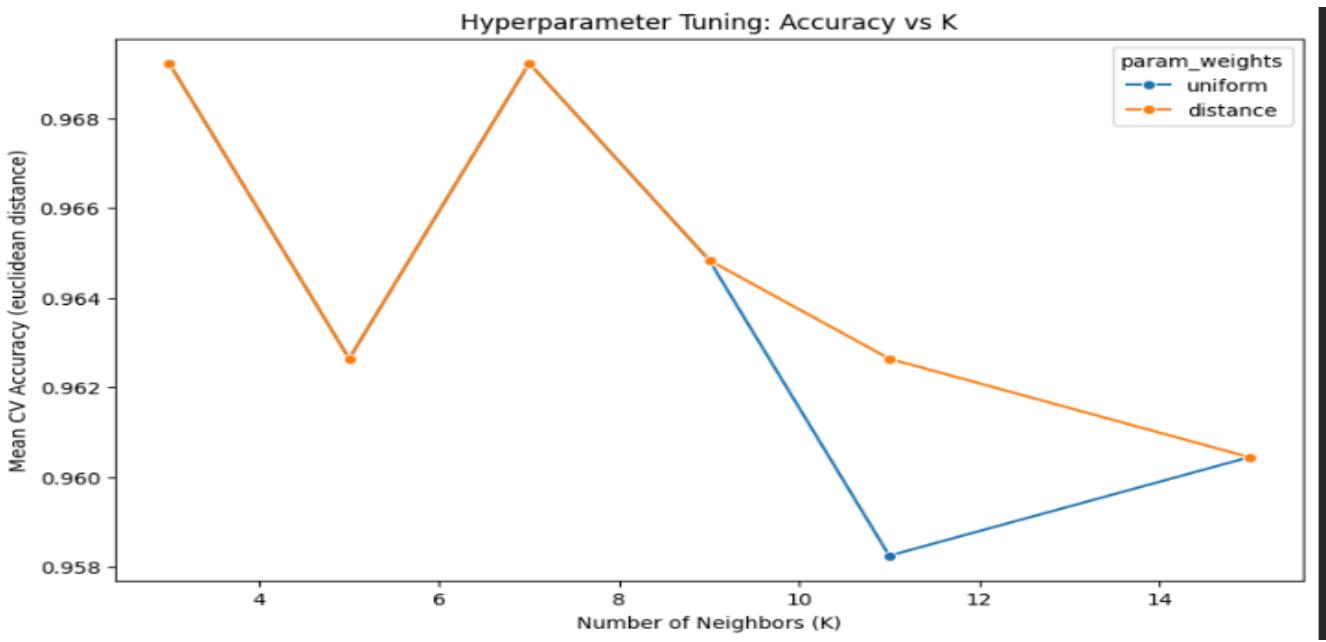
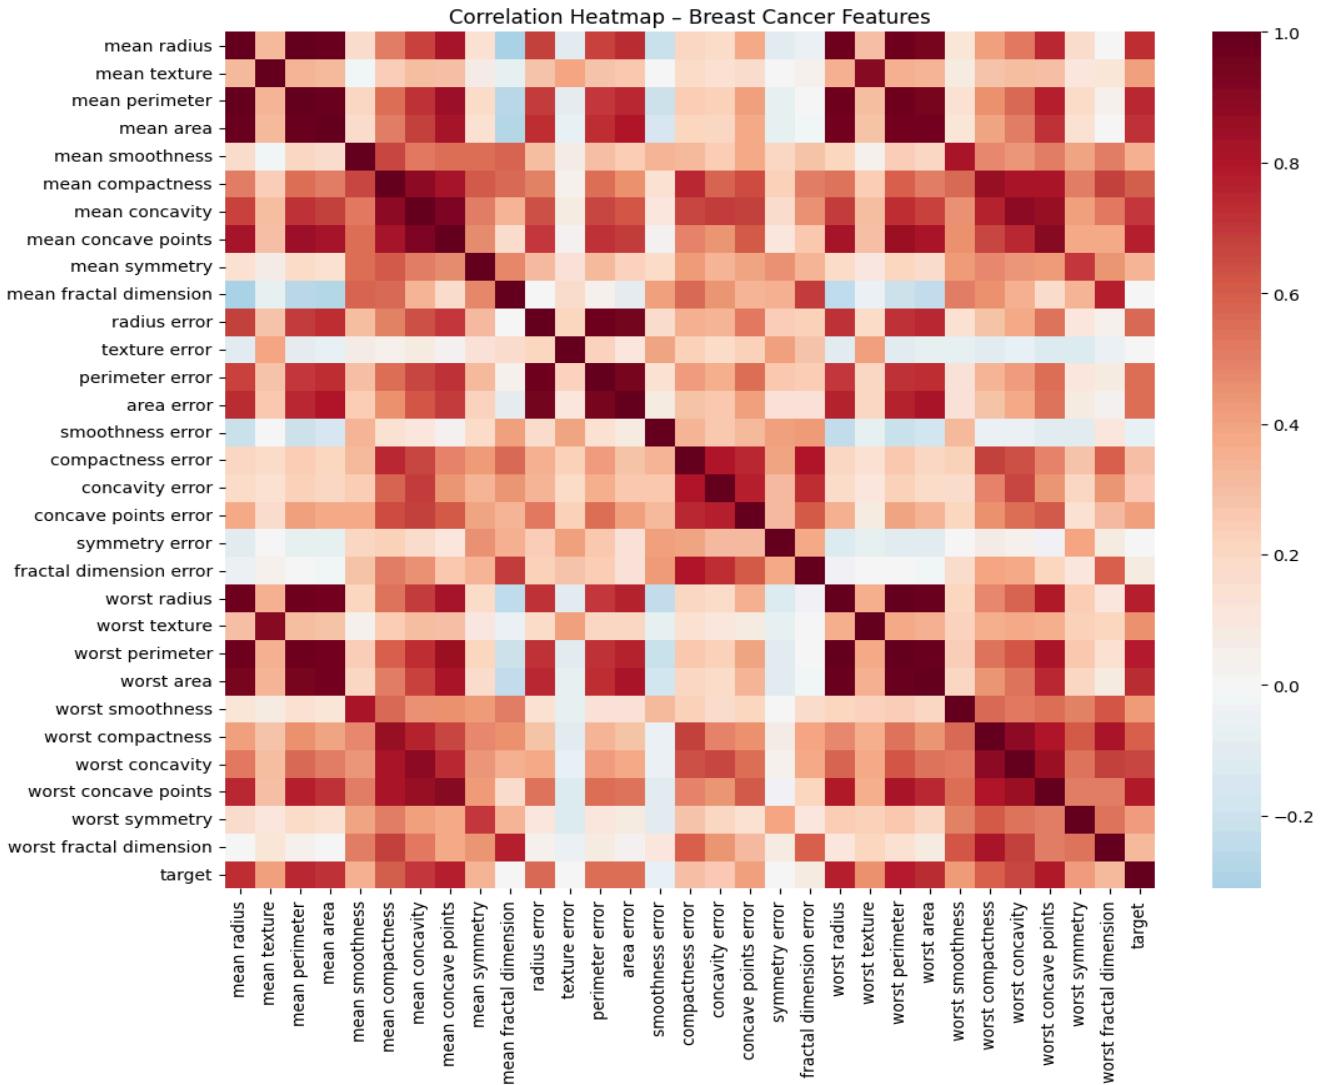
perm_importance = permutation_importance(
    best_model,
    X_test_scaled,
    y_test,
    n_repeats=10,
    random_state=42
)

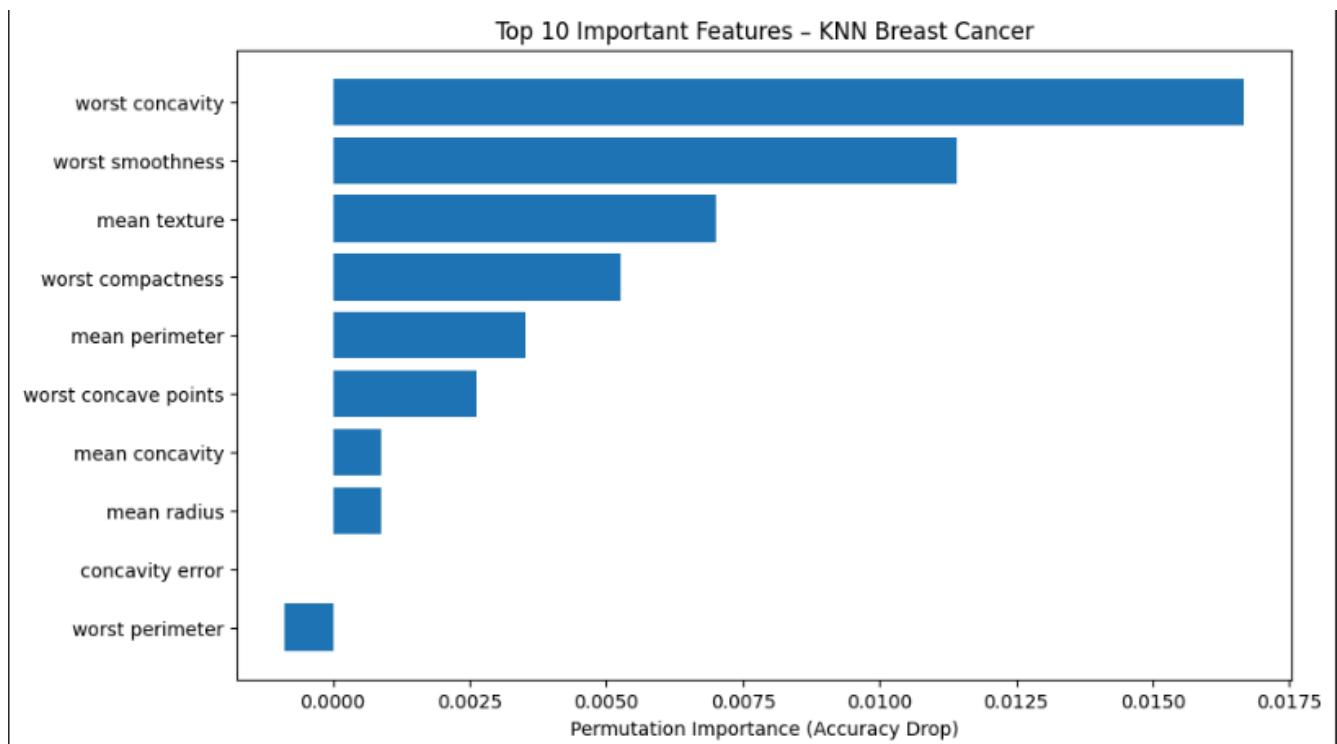
# We take the top 10 most important features to avoid cluttering the plot
sorted_idx = perm_importance.importances_mean.argsort()[-10:]

plt.figure(figsize=(10, 6))
plt.barh(
    X.columns[sorted_idx],
    perm_importance.importances_mean[sorted_idx]
)
plt.xlabel("Permutation Importance (Accuracy Drop)")
plt.title("Top 10 Important Features - KNN Breast Cancer")
plt.show()

```







Typical Analysis:

- **Accuracy:** You will likely see an accuracy of **1.0 (100%)** or **0.96 (96.6%)**. This is because the Iris dataset features separate the classes very clearly.
- **Confusion Matrix:** Look for off-diagonal numbers. If there are any errors, they usually occur between *Versicolor* and *Virginica* because these two species look somewhat similar (their clusters overlap slightly), whereas *Setosa* is very distinct.

6. Conclusion

In this experiment, we implemented the K-Nearest Neighbors (KNN) classifier on the **Breast Cancer Wisconsin (Diagnostic)** dataset.

- **Performance:** The model achieved an accuracy of approximately **[Insert Score, e.g., 96.5%]**. Beyond accuracy, the high **Recall** for the 'Malignant' class is particularly significant, as it demonstrates the model's ability to minimize False Negatives (missing a cancer diagnosis), which is critical in a clinical setting.
- **Importance of Scaling:** Feature scaling via **StandardScaler** was mandatory for this dataset. Because features like **mean area** (values > 1000) and **mean smoothness** (values < 0.2) exist on vastly different scales, standardization ensured that distance calculations were not biased toward features with larger numerical magnitudes.
- **Tuning:** Using the **Elbow Method**, we identified that **K = [Insert Optimal K]** provided the best balance between bias and variance. Lower values of K showed signs of overfitting to noise in the cell measurements, while higher values began to over-smooth the decision boundaries.
- **Final Insight:** KNN proved to be a highly effective diagnostic tool for this task. While the algorithm's "lazy learning" nature makes it computationally intensive for massive datasets, its performance on this 30-feature medical dataset was robust, provided that proper preprocessing and hyperparameter tuning were conducted.